

Deep Learning for Modeling: Concepts, Tools, and Techniques

Week 6: Momentum-based Gradient Descent and Learning Rates

Li Shuo-Hui

Norm layers

Parameter initialization principle:

- (1). Activation values should have a mean of zero**
- (2). Activation values should have the same variance across different layers**
- (3). Weights should be independent and identically distributed (i.i.d)**
- (4). Weights distribution should have a mean of zero**
- (5). Bias can be initialized to zeros**

Norm layers

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_{1...m}\}$;

Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

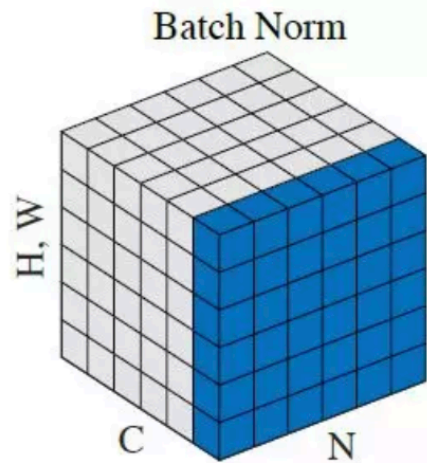
$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

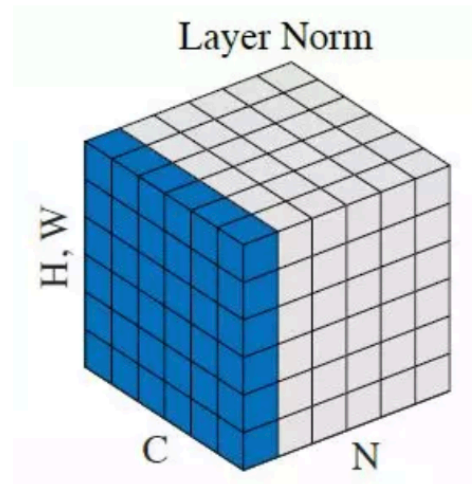
$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

Norm layers



Usually seen in CV tasks



Usually seen in NLP tasks

Gradient-based optimization: revisit

$$\min \mathcal{L}(\mathbf{x})$$

Update: $\mathbf{x}_{k+1} = \mathbf{x}_k + \eta_k \mathbf{p}_k$

η_k \rightarrow *Learning rate/step size*

\mathbf{p}_k \rightarrow *Optim direction*

$\mathbf{p}_k = -\mathbf{B}_k^{-1} \frac{\partial \mathcal{L}}{\partial \mathbf{x}_k}$

\mathbf{B}_k^{-1} \rightarrow *Local geometry*

Gradient descent: revisit

Gradient descent

Repeat

$$\boldsymbol{\theta}' = \boldsymbol{\theta} - \eta \frac{\partial \mathcal{L}_{\boldsymbol{\theta}}(\mathbf{x})}{\partial \boldsymbol{\theta}}$$

Till $\mathcal{L}_{\boldsymbol{\theta}}(\mathbf{x})$ is small enough

$$\eta_k = \eta$$

$$\mathbf{p}_k = -\frac{\partial \mathcal{L}}{\partial \boldsymbol{\theta}}$$

$$\mathbf{B}_k^{-1} = \mathbf{I}$$

Newton's method: revisit

Newton's method

Solve $\operatorname{argmin}_{\mathbf{x}} f(\mathbf{x})$

$$\frac{\partial f(\mathbf{x})}{\partial \mathbf{x}} = 0 \longrightarrow \frac{\partial f}{\partial \mathbf{x}} + \delta \mathbf{x} \frac{\partial^2 f}{\partial \mathbf{x}^2} = 0$$

Repeat $\mathbf{x}' = \mathbf{x} - \delta \mathbf{x}$ till converge

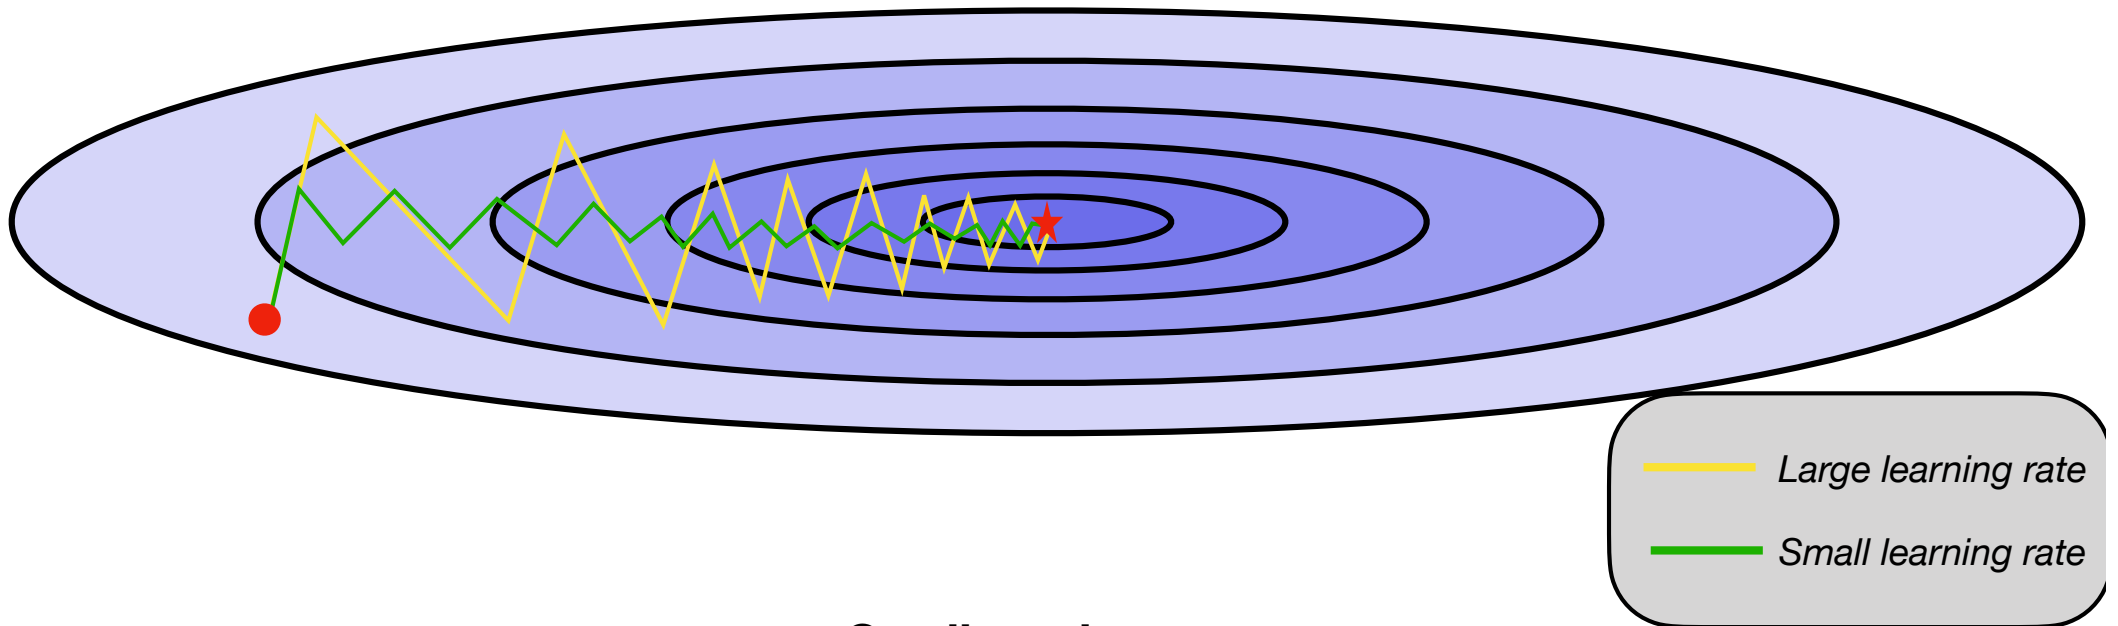
$$\eta_k = 1$$

$$\mathbf{p}_k = -\left(\frac{\partial^2 f}{\partial \mathbf{x}^2}\right)^{-1} \frac{\partial f}{\partial \mathbf{x}}$$

$$\mathbf{B}_k^{-1} = \left(\frac{\partial^2 f}{\partial \mathbf{x}^2}\right)^{-1}$$

Gradient-based optimization: problem one

The local geometry

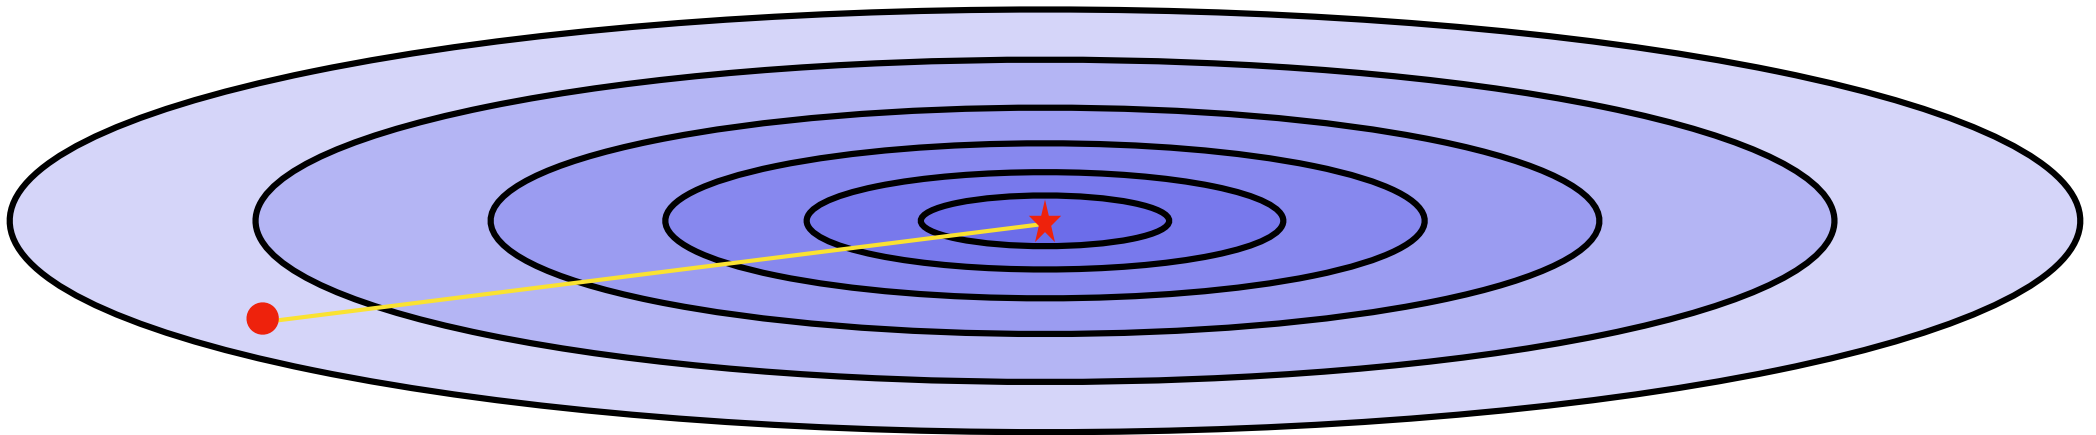


Gradient descent

$$\mathbf{B}_k^{-1} = \mathbf{I}$$

Gradient-based optimization: problem one

The local geometry

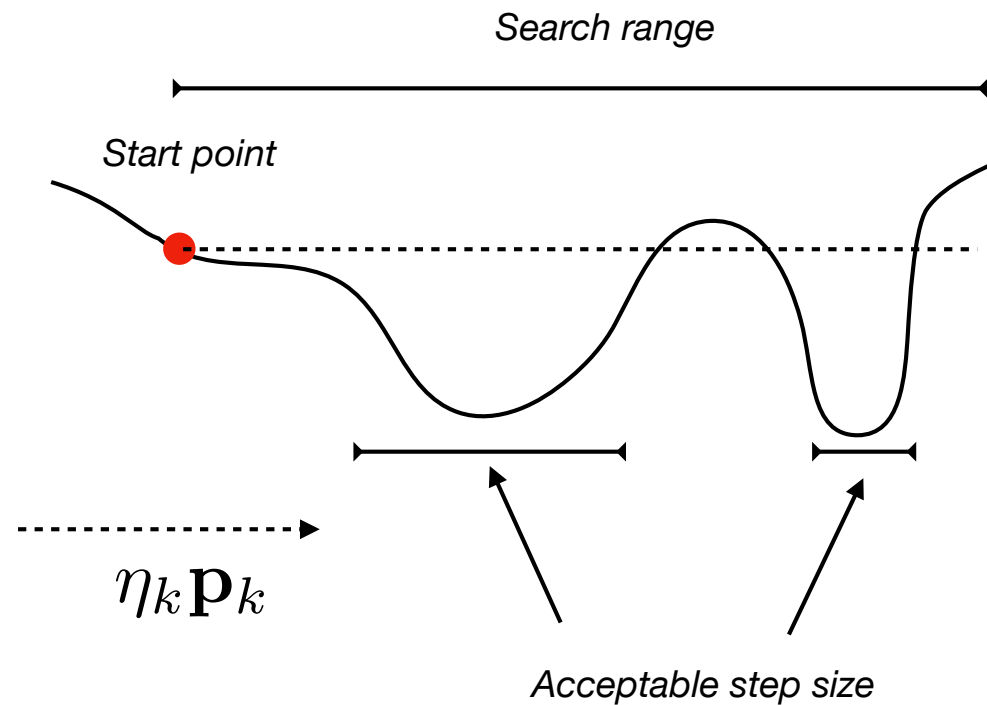


Newton's method

$$\mathbf{B}_k^{-1} = \left(\frac{\partial^2 f}{\partial \mathbf{x}^2} \right)^{-1} \quad \text{Balance update across directions!}$$

Gradient-based optimization: problem two

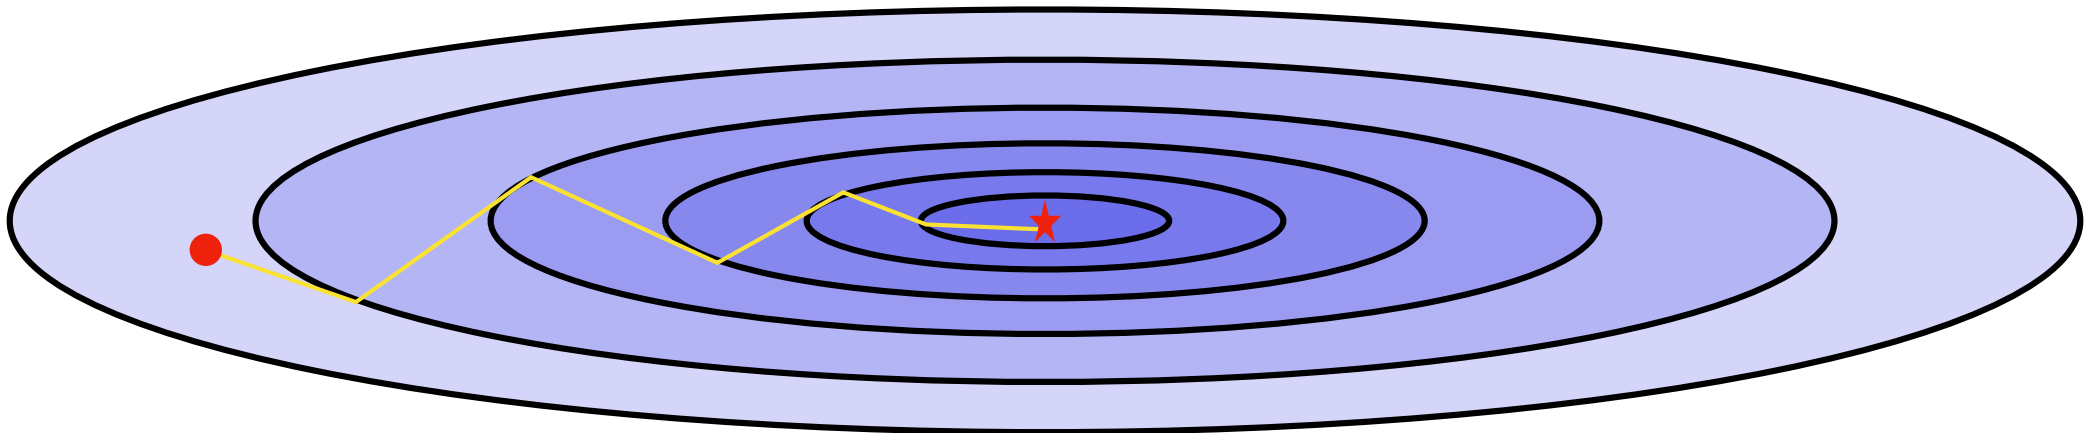
The step size selection



The line search algorithms

Gradient-based optimization: problem one

The step size selection









Steepest Descent:
gradient descent with exact line search

$$\mathbf{B}_k^{-1} = \mathbf{I} \quad \text{"Zig-zag track" tangents to level set}$$

Gradient-based optimization

Comparison

	Line search step size	Local geometry acknowledge
Newton w/ line search		
Steepest descent		
SGD		

These can be fixed with momentum and changing learning rate

Momentum-based gradient descent

Simple momentum

$$\mathbf{p}_k = -\rho \mathbf{p}_{k-1} - \frac{\partial f}{\partial \mathbf{x}_k}$$

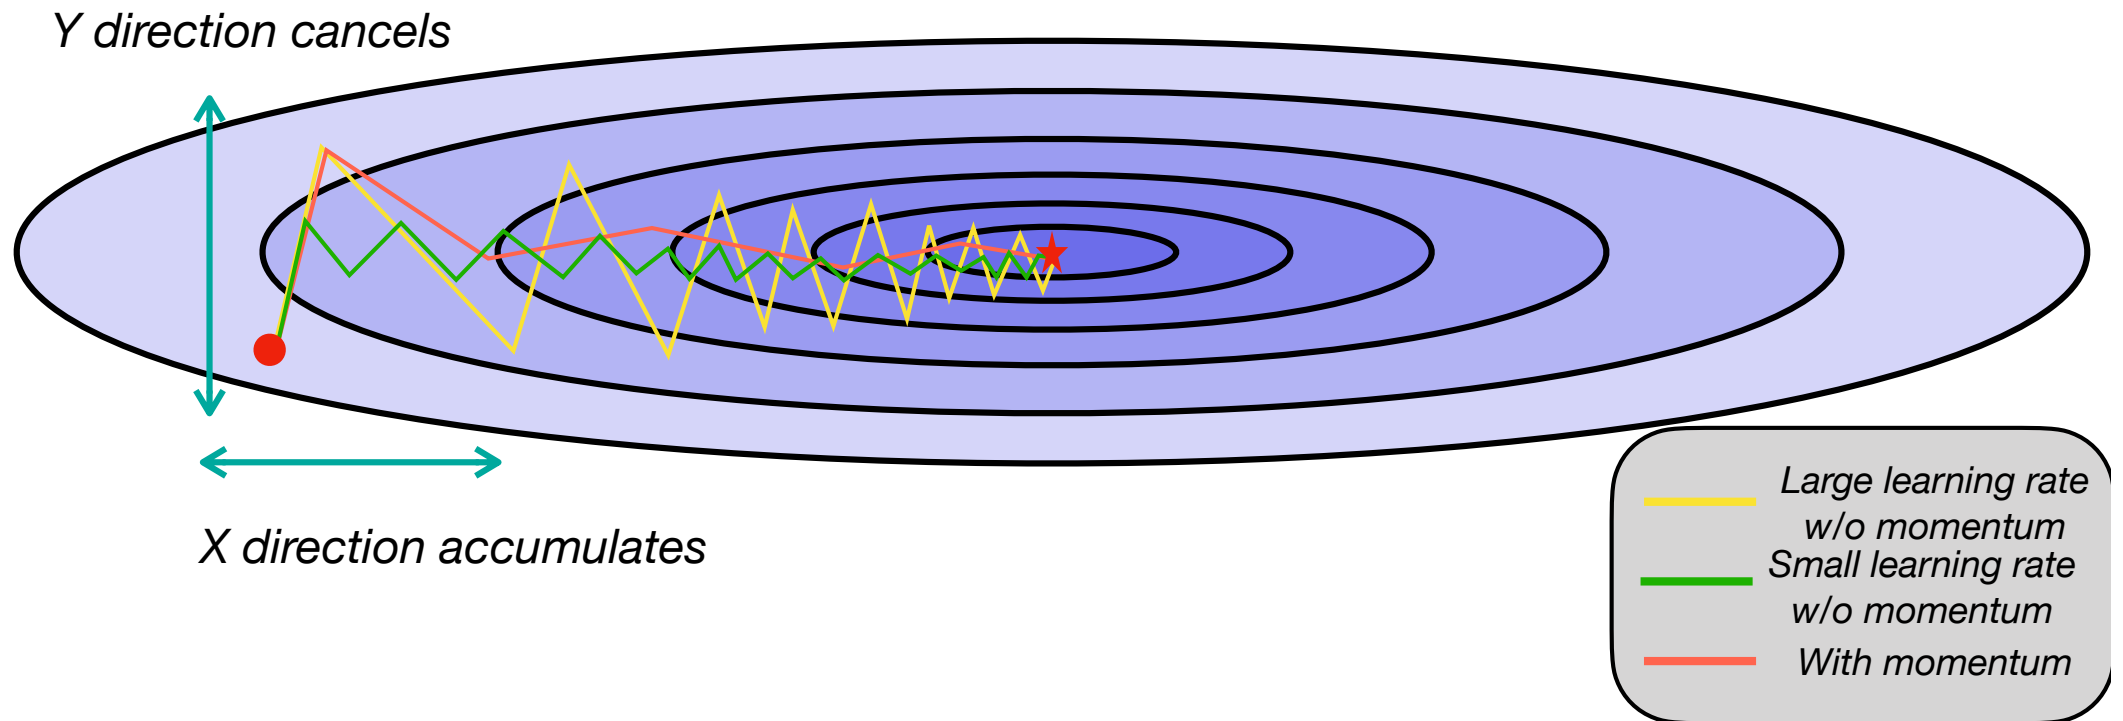
$$\mathbf{x}_{k+1} = \mathbf{x}_k + \eta \mathbf{p}_k$$

Momentum term

*Momentum decay factor
Or “friction”*

Momentum-based gradient descent

$$\mathbf{p}_k = -\rho \mathbf{p}_{k-1} - \frac{\partial f}{\partial \mathbf{x}_k}$$



Changing learning rates

Decaying learning rate

Exponential

$$\eta_k = \eta_0 \cdot \gamma^{\max(0, \lceil \frac{k-k_0}{s} \rceil)}$$

1/t scheme

$$\eta_k = \frac{\eta_0}{1 + \max(0, \lceil \frac{k-k_0}{s} \rceil)}$$

⋮

Decaying learning rate to adapt finer optimization at the later stage

Changing learning rates

Adaptive learning rate

Adagrad

$$\eta_{k,i} = \frac{\eta_0}{\sqrt{\sum_j^k p_{j,i}^2}}$$

$$\begin{pmatrix} x_{k+1,1} \\ x_{k+1,2} \\ \vdots \\ x_{k+1,n} \end{pmatrix} = \begin{pmatrix} x_{k,1} \\ x_{k,2} \\ \vdots \\ x_{k,n} \end{pmatrix} + \begin{pmatrix} \frac{\eta_0}{\sqrt{\sum_j^k p_{j,1}^2}} \\ \frac{\eta_0}{\sqrt{\sum_j^k p_{j,2}^2}} \\ \vdots \\ \frac{\eta_0}{\sqrt{\sum_j^k p_{j,n}^2}} \end{pmatrix} \odot \begin{pmatrix} p_{k,1} \\ p_{k,2} \\ \vdots \\ p_{k,n} \end{pmatrix}$$

Adaptively tune the learning rate per parameter using historical updates

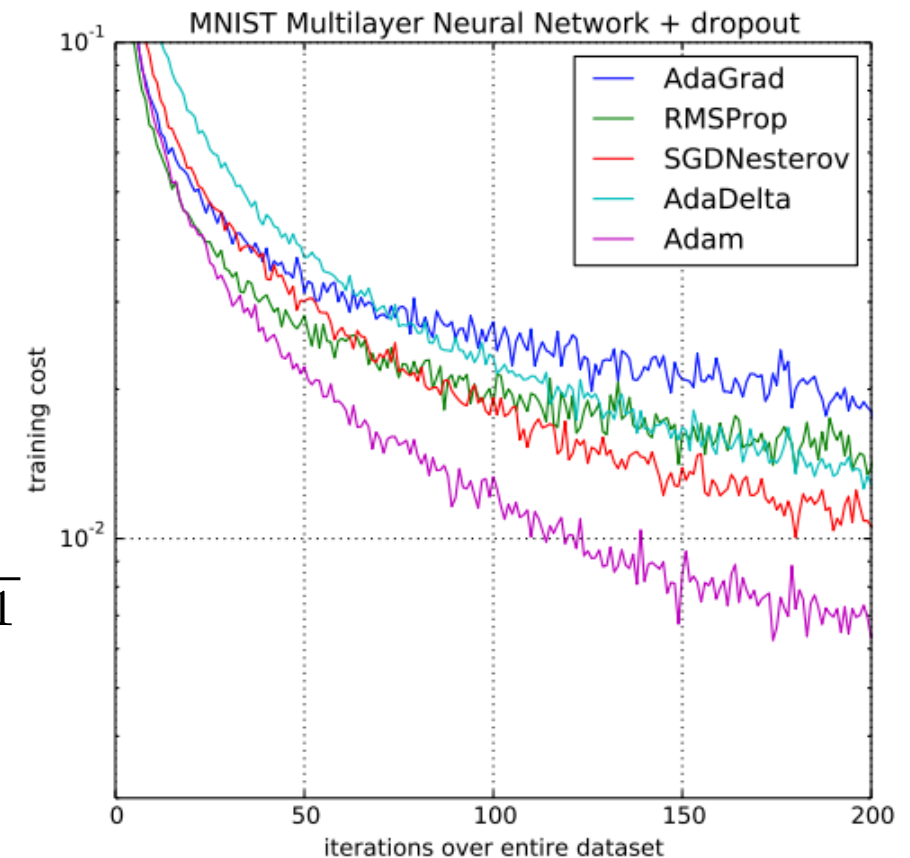
Adam optimizer

$$m_{k+1} = \beta_1 m_k + (1 - \beta_1) \frac{\partial \mathcal{L}}{\partial \theta_k}$$

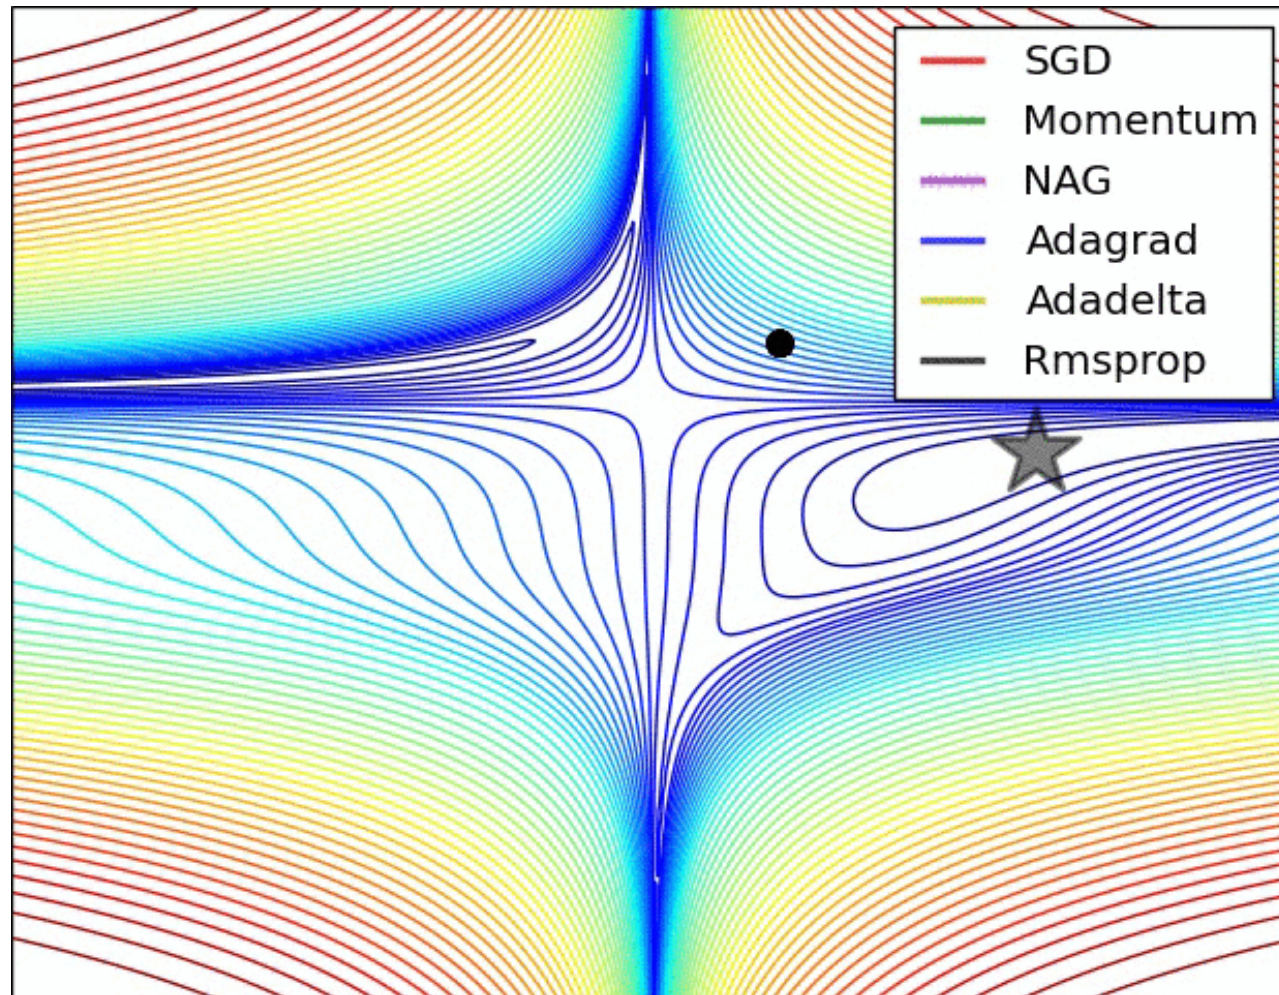
$$v_{k+1} = \beta_2 v_k + (1 - \beta_2) \frac{\partial \mathcal{L}}{\partial \theta_k}^2$$

$$\hat{m}_{k+1} = \frac{m_{k+1}}{1 - \beta_1^{k+1}} \quad \hat{v}_{k+1} = \frac{v_{k+1}}{1 - \beta_2^{k+1}}$$

$$\theta_{k+1} = \theta_k - \eta \frac{\hat{m}_{k+1}}{\sqrt{\hat{v}_{k+1}}}$$



Some intuitions



Deep Learning for Modeling: Concepts, Tools, and Techniques

Week 6 Tutorial

Li Shuo-Hui

PyTorch Module

nn.Module

```
NeuralNetwork(  
    (hidden): Linear(in_features=10, out_features=30, bias=True)  
    (sigmoid): Sigmoid()  
    (output): Linear(in_features=30, out_features=2, bias=True)  
    (softmax): Softmax()  
)
```

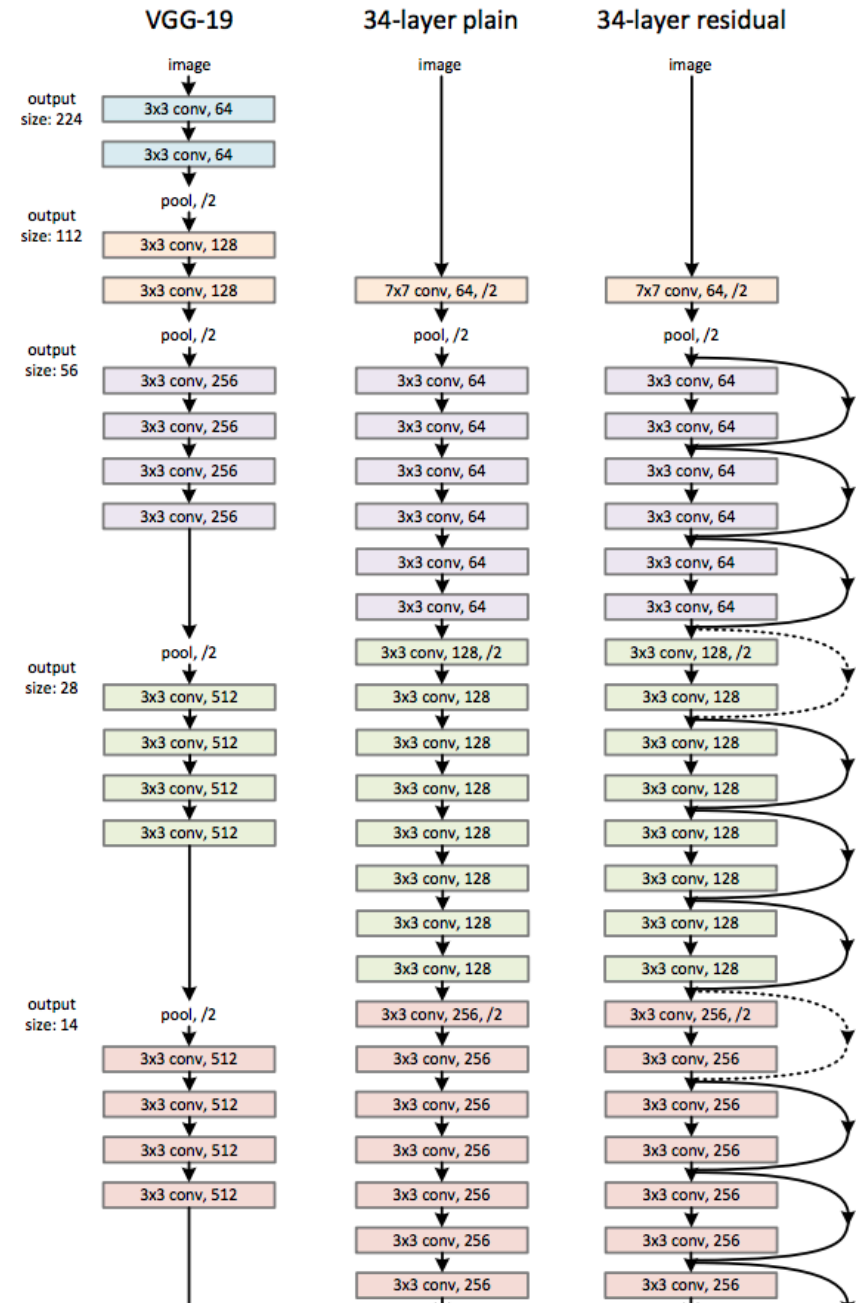
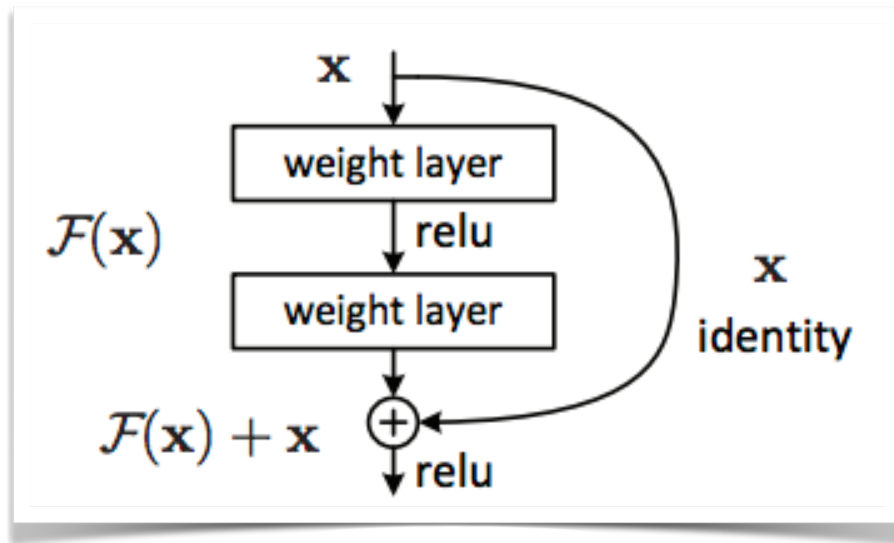
nn.Sequential

```
Sequential(  
    (0): Linear(in_features=10, out_features=30, bias=True)  
    (1): Sigmoid()  
    (2): Linear(in_features=30, out_features=2, bias=True)  
    (3): Softmax()  
)
```

nn.ModuleList

```
NeuralNetwork(  
    (module_list): ModuleList(  
      (0): Linear(in_features=10, out_features=30, bias=True)  
      (1): Sigmoid()  
      (2): Linear(in_features=30, out_features=2, bias=True)  
      (3): Softmax()  
    )  
)
```

Resnet



Unet

